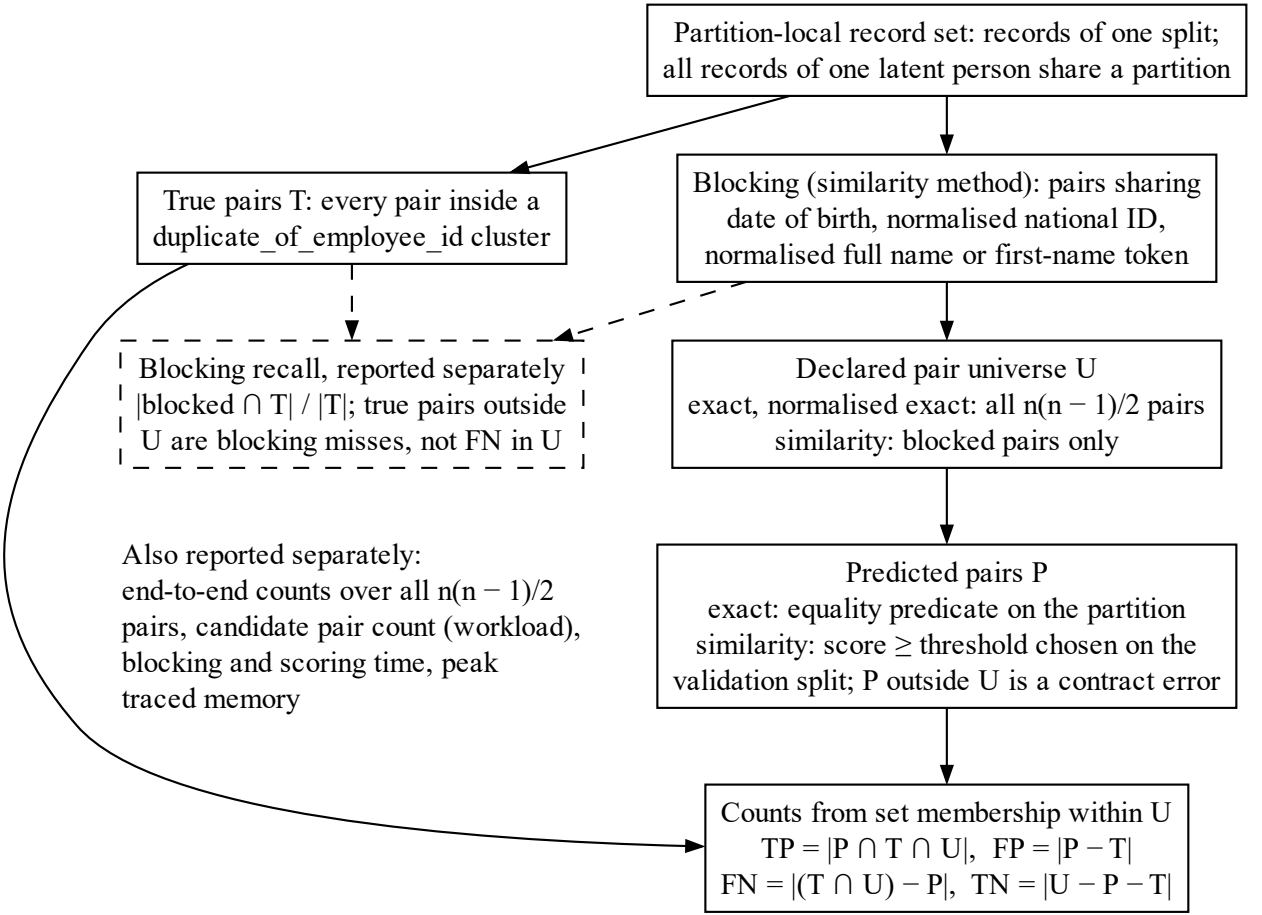
Partition-local record set: records of one split; all records of one latent person share a partition
True pairs T: every pair inside a duplicate_of_employee_id cluster
Blocking (similarity method): pairs sharing date of birth, normalised national ID, normalised full name or first-name token
Blocking recall, reported separately |blocked ∩ T| / |T|; true pairs outside U are blocking misses, not FN in U
Declared pair universe U
exact, normalised exact: all n(n − 1)/2 pairs
similarity: blocked pairs only
Also reported separately:
end-to-end counts over all n(n − 1)/2 pairs, candidate pair count (workload), blocking and scoring time, peak traced memory
Predicted pairs P
exact: equality predicate on the partition
similarity: score ≥ threshold chosen on the validation split; P outside U is a contract error
Counts from set membership within U
TP = |P ∩ T ∩ U|, FP = |P − T|
FN = |(T ∩ U) − P|, TN = |U − P − T|